

Big Data Technologies

Agenda

- Apache Hadoop
 - MapReduce
 - YARN

Hadoop

Map-Reduce Programming

MR Job XML Configuration

- `fs.defaultFS = hdfs://namenode:9000/`
- `mapreduce.framework.name = yarn`
- `yarn.resourcemanager.address = resourcemanager:8032`
- Other config parameters can be set up as per application requirement e.g. `dfs.replication`, `dfs.blocksize`, or additional user defined parameters.
- XML configuration is given using "-conf" option.

Understanding MR execution

- MR summary (on console)
 - Number of mapper & reducer tasks.
 - Number of input & output records for mapper
 - Number of input & output records for reducer
- MR log files review
 - `$HADOOP_HOME/logs`
 - MR jobs logs --> `$HADOOP_HOME/logs/userlogs`
 - The MR job logs --> `$HADOOP_HOME/logs/userlogs/application_appid`
 - The MR mapper/reducer logs --> `$HADOOP_HOME/logs/userlogs/application_appid/container_id`

- stdout/stderr
- syslog

Custom Job counters

- Hadoop Counters provides a way to measure the progress or the number of operations that occur within map/reduce job.
- Counters in Hadoop MapReduce are a useful way for gathering statistics about the MapReduce job for quality control or for application-level.
- Example:
 - step 1: Create enum of custom Job counter

```
enum NcdcCounters {  
    BAD_RECORDS, BAD_READINGS, GOOD_READINGS  
}
```

- step 2: Access the counter in Mapper or Reducer program and increment it.

```
Counter cntr = context.getCounter(NcdcCounters.GOOD_READINGS);  
cntr.increment(1);
```

Input/Output Format and Input Splits

- InputFormat – how to read data
 - FileInputFormat, TextInputFormat
 - KeyValueTextInputFormat, NLineInputFormat
 - DBInputFormat
 - CombineTextInputFormat
- RecordReader (nested class in InputFormat) – Logical division of record

- Number of mappers = Number of input splits
- Number of input splits \approx Number of HDFS blocks
- Typically mapper tasks are executed on the system on which data blocks (or replica) are present. Such mapper tasks are called as "Data-local Map tasks".
- OutputFormat – how to write data
 - FileOutputFormat, TextOutputFormat
 - DBOutputFormat
- RecordWriter (nested class in OutputFormat) – Write individual record
- Number of output files = Number of reducers
- Output written on HDFS (replicated)

Data flow of MR job (Single reducer)

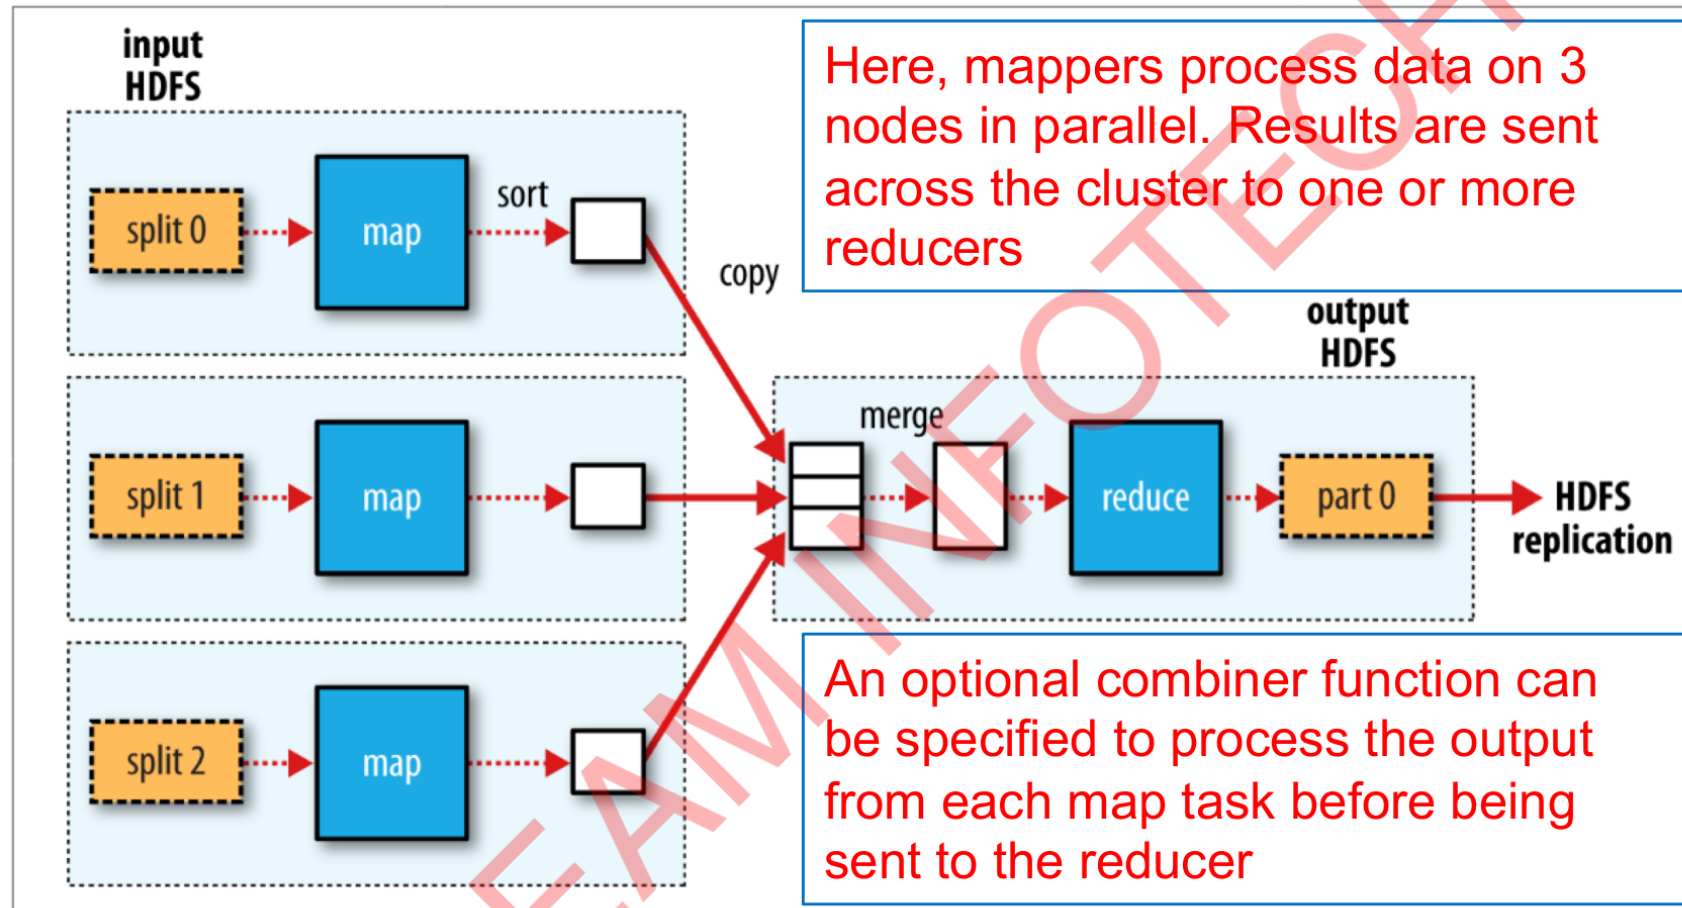


Figure 2-3. MapReduce data flow with a single reduce task

Combiner

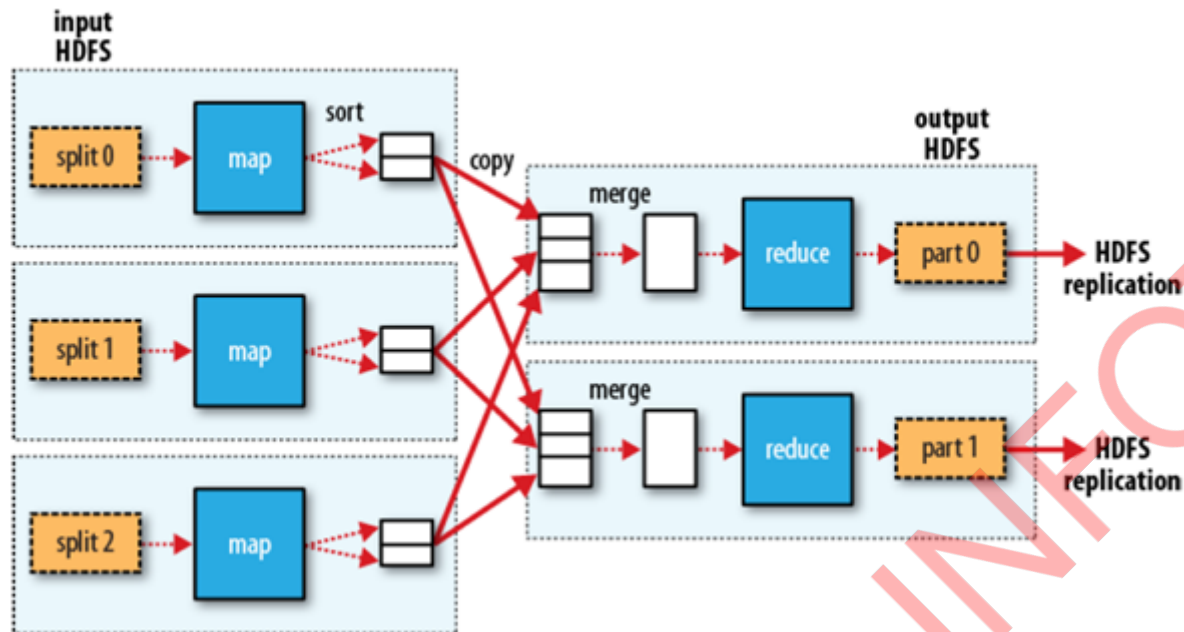
- Combiner is a local reducer i.e. runs reducer (aggregation logic) within mapper task process.
 - Minimize output of mapper task
 - Less merge & shuffle
 - Less network transfer

- Less aggregation in reducer
- Combiner is optional.
- Works for commutative & associative aggregate functions only.
 - $A + B = B + A$
 - $A + (B + C) = (A + B) + C$

Partitioner

- By default MR job have single reducer.
- Having huge data for aggregation may lead to out of memory error.
- Number of reducers can be configured in job configuration file OR in driver code.
 - `job.setNumReduceTasks(2);`
 - `mapreduce.job.reduces = 2`
- Number of partitions = Number of reducers
- Output of mapper is divided into multiple partitions based produced key.
- By default HashPartitioner is used, that distributes mapper output in number of partitions uniformly. Refer slides.
- Can implement custom partitioner class (inherited from Partitioner class), if HashPartitioner is not suitable (not balance load on reducers).
 - `job.setPartitionerClass(CustomPartitioner.class);`

Data flow of MR job (Multiple reducer)



Hadoop MR data flow (detailed)

- Mapper Task/Process
 - HDFS --> InputFormat+RecordReader --> Mapper --> Partitioner --> Ring Buffer (100MB) --> Sort & Group --> Combiner --> Write on Disk.
- Reducer Task/Process
 - Mapper Task Output --> Reducer input buffer --> Merge Sort & Group --> Reducer --> OutputFormat+RecordWriter --> HDFS.
- Refer slides
- <https://0x0fff.com/hadoop-mapreduce-comprehensive-description/>

Map-only Jobs

- If data processing includes only data cleaning/filtering (per record) and no aggregation operations, we create Map-only job.
- Map only job is created by setting number of reducers to zero.
 - `job.setNumReduceTasks(0);` // Java code -- OR
 - `mapreduce.job.reduces = 0` // Java xml config file

- Applications
 - Data ingestion (from different sources)
 - Data cleansing
 - Data filtering
 - Data pre-processing
- SQOOP is Hadoop eco-system/application that generates Map-only jobs for transferring data from RDBMS to Hadoop and vice-versa.
 - RDBMS --> sqoop import --> Hadoop
 - Hadoop --> sqoop export --> RDBMS

Reduce-only Jobs

- It is not possible to have Reduce only jobs.
- However, we can implement Map-Reduce jobs where mapper is not doing any significant processing. The output of mapper will be same as input data. Such mapper is called as IdentityMapper.
- The output of IdentityMapper will be sorted (by null key) and sent to the reducer. The reducer code should be written accordingly.
- This is not commonly used.