

Cassandra

column based NoSQL database




Introduction


@2003 => GFS -> Google File System => Distributed FS }
HDFS => hadoop Distributed file system => HBase

- Google BigTable ✓
 - High performance data storage system built on GFS and other Google technologies
 - Master-slave architecture
 - One key, multiple values
 - Columnar, SSTable (Sorted String Table) Storage, Append-only, Memtable, Compaction
- Amazon DynamoDb ✓
 - Highly available and scalable key-value storage system => horizontal scaling
 - Decentralized peer to peer architecture
 - Compromise on consistency for better availability - Eventual consistency
 - Consistent hashing, Gossip protocol, Replication, Read repair
- Cassandra ✓
 - Inherited from BigTable and DynamoDb
 - BigTable: Column families, Memtable, SSTable
 - DynamoDb: Consistent hashing, Partitioning, Replication ,

RAM



SPOT



	ename	job	sal
1	abc	1 sales	1 2000
2	xyz	2 MGR	2 3000
3	pqr	3 Dev	3 2500

Arrows from the row numbers (1, 2, 3) in the table point to the **Row Key** label.



Introduction

- **Developed by**

- Avinash Laxman (Co-inventor Amazon DynamoDb)
- Prashant Malik (Technical Leader at Facebook)

- **Goals:**

- ✓ Distributed NoSQL database (on commodity hardware)
- ✓ Large amount of structured data
- ✓ High availability
- ✓ No single point of failure

- Basic data model is rows & columns
- Column-oriented, Decentralized peer to peer & follow Eventual consistency
- Datastax company develop and support commercial edition of Cassandra



Cassandra Development

- Developed in Java
- 2007-2008 - Developed at Facebook
- July 2008 - Open sourced by Facebook
- March 2009 - Apache Incubator project
- February 2010 - Apache Top-level project
- 2011 - version 0.8 - Added CQL
- 2013 - version 2.0 - Added light-weight transactions, Triggers
- 2015 - version 3.0 - Storage engine improved, Materialized views
- 2017 -version 3.11 – bug fix from the last release ✓
- 2021 - version 4.0.5 ---> we r using this version
- 2022 - version 4.1.0 - Latest release



Installation

- Prerequisite
 - Java 8 (~~Java 11~~ experimental)
- Can be installed through apt or yum tool (Ubuntu/CentOS)
- Manual installation
 - Download Cassandra 3.11.x (.tar.gz) and extract it
 - set CASSANDRA_HOME to Cassandra directory
 - set JAVA_HOME to JDK 8 directory
 - Install python 2.7 (for cqlsh)
 - set CASSANDRA_HOME/bin into PATH variable
 - Start Cassandra
- terminal1> cassandra
- terminal2> cqlsh

→ mapping

mysql	cassandra	
-----	-----	
database	keyspace	✓
table →	table	
row →	row	
column →	column	



Features

- Peer to peer architecture → no master-slave => single point of failure possibility
- Linear scale performance -> capacity \propto peers
- High Performance → high speed read write
- Simplified deployment and maintenance => on linux
- Less expensive horizontal scaling
- Supports multiple programming languages => java, python, ruby, C#, REST services
- Operational and Development simplicity => CQL similar to SQL
- Cloud Availability AWS, Azure, GCP
- Ability to deploy across data centers
- Fault tolerant
- Configurable consistency (tight or eventual)
- Flexible data model
- Column family store → number of columns can be changed dynamically



Limitations

- Aggregation operations are not supported
- Range queries on partition key are not supported
- Not good for too many joins
- Not suitable for transactional data
- During compaction performance / throughput slows down
- Not designed for update-delete



Performance

- Performance measures
 - Throughput (operations per second) (high)
 - Latency time required for single operation (low)
- Cassandra vs MySQL
 - MySQL (more than 50GB data)
 - Write speed: 300 ms
 - Read speed: 350 ms
 - Cassandra (more than 50GB data)
 - Write speed: 0.12 ms
 - Read speed: 15 ms



Applications

- Applications
 - Product catalog/Playlist
 - Recommendation/Personalization engine
 - Sensor/IoT data
 - Messaging/Time-series data
 - Fraud detection
- Customers
 - Facebook, Netflix, eBay, Apple, Walmart, GoDaddy ,
- Application requirements
 - Store and handle time-series data
 - Store and handle large volume of data
 - Scale predictably (Linear Scaling)
 - High availability



Data Model

- Cassandra provides the Cassandra Query Language (CQL), an SQL-like language, to create and update database schema and access data
- CQL allows users to organize data within a cluster of Cassandra nodes using:
 - **Keyspace** ✓
 - Defines how a dataset is replicated, per datacenter
 - Replication is the number of copies saved per cluster. Keyspaces contain tables
 - **Table**
 - Defines the typed schema for a collection of partitions. Tables contain partitions, which contain rows, which contain columns. Cassandra tables can flexibly add new columns to tables with zero downtime.
 - **Partition**
 - Defines the mandatory part of the primary key all rows in Cassandra must have to identify the node in a cluster where the row is stored
 - All performant queries supply the partition key in the query.
 - **Row**
 - Contains a collection of columns identified by a unique primary key made up of the partition key and optionally additional clustering keys
 - **Column**
 - A single datum with a type which belongs to a row



- Users can access Cassandra through its nodes using Cassandra Query Language (CQL)
- CQL treats the database (Keyspace) as a container of tables
- Programmers use cqlsh: a prompt to work with CQL or separate application language drivers



- ascii: US-ascii character string
- bigint: 64-bit signed long ints
- blob: Arbitrary bytes in hexadecimal
- boolean: True or False
- counter: Distributed counter values 64 bit
- decimal: Variable precision decimal
- double: 64-bit floating point
- float: 32-bit floating point user defined Datatype
- frozen: Tuples, collections, UDT containing CQL types
- inet - IP address in ipv4 or ipv6 string format
- int: 32 bit signed integer
- list: Collection 01 elements
- map: JSON style collection of elements
- set: Sorted collection of elements
- text: UTF-8 encoded strings
- timestamp: ID generated with
- date+time: as int/string
- timeuuid: Type 1 uuid
- tuple: A group of 2,3 fields
- uuid: Standard uuid (128-bit)
- varchar: UTF-8 encoded string
- varint: Arbitrary precision integer



Cassandra vs MongoDB: Differences

■ Cassandra

- Dev In java ✓
- Column based ✓
- Cassandra uses a traditional model with a table structure, using rows and columns.
- Cassandra offers an assortment of peers node
- Used in-house query language, CQL
- no internal aggregation framework
- According to the CAP theorem Cassandra is an AP system.

■ MongoDB

- Dev in C++ ✓
- Document based ✓
- MongoDB employs an objective-oriented or data-oriented model
- MongoDB uses a single master node.
- queries are structured into JSON fragments
- own aggregation framework
- According to the CAP theorem, MongoDB is a CP system

